

Modelling and model checking interactive programs

Here we are interested in programs that interact with the world around them and need not terminate to be useful. These programs run concurrently with the world around them. The formal modelling of such programs has been achieved in a wide variety of styles. Unfortunately each style of modelling take a significant effort to understand in detail and consequently most people become an *expert* in only one way of modelling. Each of the modelling styles is predicated on a set of assumptions that are commonly unvoiced and hence comparison between styles is far from easy.

To help formalise this we model programs that can themselves be concurrent. The class of such concurrent programs is very varied and many distinct styles of analysis appear in the literature. We find it helpful to decompose these concurrent programs into two broad, overlapping, classes *State based* and *Event based*.

Here we formally model interactive processes by first formally capturing their operational behaviour and then defining process equivalence on the operational semantics. Equality has the following properties:

1. it is a reflexive, symmetric and transitive relation,
2. equal processes can be substituted and
3. equal processes can not be distinguished.

We will use automata, Labeled Transition Systems LTS, to model the operational behaviour of processes and define operator on automata to compose processes and there by building larger automata. From our perspective the most important operator is parallel composition as we use this define how processes are "distinguished" and hence we use parallel composition to define process equality. To do this we, later, will formalise Testing in a way that can be applied in many situations.

By concurrent state based programs interact by sharing memory and concurrent event based programs interact by sharing events. Our event based programs are commonly referred to as *processes*. On the one hand the same real world object can be formalised in a state based or event based fashion and on the other hand different languages implement shared memory of shared events as primitive.

We believe that state based models can be implemented on event based primitives and event based models can be implemented on state based primitives. The usefulness of having different styles to formalise the same thing is that different objects are easier to formalise with each of the styles.

1 A practical approach to formalisation

Theory has an internal virtue in supporting the exploration of the consequences of a small set of assumptions. In addition it gives confidence in adopting analytic methods in a practical setting. Nonetheless their are circumstances where we are not interested in a *theoretical ideal* if our goal is to provide sound usable tools for the analysis of systems.

Our primary style of modelling can be described as based on *hand shake events* and as is common we call such programs processes. We consider CSP, CCS and ACP as all using hand shake events and like here make the following assumptions:

1. we abstract away details about time

2. an observable event in one process can require the existence of another process to be ready to perform a *related event* before it is executed. This results in a branching structure and would require the use of branching temporal logic if one decided to use logic to specify the details of a processes behaviour.
3. an event is either *observable and blockable* or *unobservable and unblockable* (Note the DES community separate out these properties and hence have four types of event)
4. we abstract away details about one event *causing* a related event to occur. Consequently all events in a set of related events are treated equally (Note modelling causality between synchronising events has been shown to be useful when relating handshake events with broadcast events)
5. a process can be in a state where it is prepared to perform any event from some set of events. How this is achieved is not specified but is a programming language primitive in Ada rendezvous, Occam events and Go events.
6. we adopt the interleaving assumption - parallel processes can be thought equivalent to sequential process

Although it is interesting to have a sound and complete inference system to reason about programs it should be borne in mind that there is no one universally agreed upon notion of equality between interacting systems and hence putting lot of effort into constructing complete inference system may be of little practical use. What we believe to be of greater practical importance than completeness is that the inference system is

1. *sound*
2. *easy to understand*
3. *push button feedback*
4. *extensible*

Both B and Event-B use of forward simulation and not backward simulation hence are inherently incomplete with respect to refinement and yet can be of practical use.

Here we are less concerned with representing systems exclusively as process terms and constructing a sound and complete set of axioms than in defining sound algorithms for the rewriting of the operational semantics of processes. This can be seen as what in the literature is called *Visual verification*.

1.1 Approaches

The **B and Event-B approach** is to start with an abstract specification of the System as a whole. This can be stepwise refined until sufficient detail has been provided and then the detailed System can be decomposed into the *Process* of interest plus the *Context* it has been designed to work in.

Two practical aspects of this approach are:

1. the abstract and more detailed concrete models are built by hand, one is not constructed from the other (not correct by construction), then the relation between the two models is verified. Thus the method can be applied in both directions, analysis by refinement and by abstraction.

2. The models defined at each level can be deterministic. The retrieve relation between the two models introduces non determinism that is removed by verifying the relation is a refinement.

This approach could be described as supporting hierarchical specifications. That is each level in the hierarchy consists of one or more specification given at the same level of detail whereas between levels the specifications are of the same thing but given at differing levels of detail. Frequently an abstract specification can be constructed from a more detailed specification by the abstraction of some detail. In our simple models that contain only atomic events this means the abstraction of some of the events.

A **Control theory approach** used quite widely, is to model the context in which the process finds its self in addition to the process itself. In this approach we have separate models for the *Process* itself and the *Context* in which it resides. The combination of the two is called the *System*.

Frequently for processes we wish to build it is easier to decompose the specifications into two steps:

1. specify the behaviour of the context in which the program will run
2. specify the behaviour of the system as a whole, that is specify the behaviour of the program and the context running together

Because we have parallel composition as a first class citizen we can further refine the Process if we need to.

Private public To build understandable models they need to be small. We show how to *abstract* the private component while preserving the public behaviour.

Finite state approximations for pragmatic reasons we favour fast push button analysis over techniques that require a high degree of human time and expertise.

Visual verification Specifying properties of interacting systems is hard and in practice engineers are reluctant to spend a lot of time learning languages that may not be helpful.

Operational semantics provides a simple mapping between state based and event based systems.

Testing semantics provides confidence in the situations for which our definitions of equality and refinement are safe to use.

Robust semantics People frequently think in different ways about the same system depending upon the task at hand. Formalising our semantics in a variety of styles helps do this with some confidence.

Extensions via theory morphisms Each of our formalisms admit a logical interpretation where refinement is implication. The construction of Galois connections between these formalisms offers us a great deal of flexibility.

1. selecting different formalisms at different point in the development
2. mixing different formalisms
3. using the relational semantics we have silent out side of frame

Look to practice When things get tough to formalise look to your engineering intuitions. Refactor your theory rather than bolt on a patch.

2 State based interaction

Concurrent programs that interact by sharing variables have been formalised in many distinct ways. Here guarded actions similarly to both TLA+ and Event B. This has the benefit of making the connection event based interaction, as formalised by process algebras, easier to formalise.

Our state based programs will be called **modules** and contain a set of named, *guarded actions*. The module contains a set of variables that each action may reference.

3 Event based interaction

We model two styles of event based interaction, the first is based on the hand shake synchronisation of process algebras and the second is based on the broadcast synchronisation of IOA automata. The syntax used to index processes is taken from LTSA.

4 Defining Basic Processes

We are interested in defining processes that have private state and only interact with other processes via **events**. Process calculi CSP, CCS, ACP and Lotos are successful examples of such event based formalisms. These event based formalisms make use of a wide range of denotational semantics and hence a wide range of process equivalences. Each can be given an axiomatic semantics to enable reasoning about the processes. We take a non axiomatic approach that loosely follows the notion of *Visual Verification* from Valmaries .. [].

Event based formalisms define processes that interact with processes running concurrently with them. Only events can be seen, interacted with, by other processes as such these observable events can be thought of as *half events* that is they need the other half to be ready before they can be executed.

Our approach requires the distinction between what is in the interface of a process and what is private to the process. Subsequently we define how to remove (*abstract away*) what is private while preserving what can be seen of the process behaviour.

In the tool we introduce process names will start with an upper case letter whereas event names will start with a lower case letter.

Sequential processes are defined using a simple process algebra $\Sigma = \{Act, -, >, |, STOP\}$ the operational semantics of the defined process definition will be rendered as an automata.

An automata is a tuple $A \triangleq (N_A, S_A, T_A, \alpha_A)$ where:

N_A is the set of automata nodes,

S_A is the set of start nodes $S_A \subseteq N_A$

T_A is the set of transitions, triples (n, a, m) where $a \in \alpha_A$ and $\{n, m\} \subseteq N_A$

α_A is the alphabet of the process, the set of events in the processes $\{a : n \xrightarrow{a} m\} \subseteq \alpha_A$

□


Processes definitions are enclosed in **processes** $\{ \dots \}$ as shown in the examples below. We will frequently omit the **processes** $\{ \dots \}$. The processes that are to be displayed as automata appear as a list **automata** $A, Ap..$

A simplest process is **STOP** the process that dose nothing. The more interesting but very basic processes that we discuss consist of a finite state space and transitions labelled with atomic events.

4.1 Event prefixing


A simple process that performs a single event and stops can be built by prefixing an event **takeTea** to the **STOP** process using the \rightarrow operator by the command:

```
automata {
  Simple = (takeTea->STOP).
}
```



Every process we define can be represented by a transition labeled automata. The events, like **takeTea**, have an informal meaning (semantics) given by relating them to some real world event. We can prefix a second event

```
Two = (teaButton->takeTea->STOP).
Two = (teaButton->Simple).
```



The two definitions of **Two** produce the same automata. If you only want to see **Two** then you can suppress the production of the initial automata **Simple** by add ***** as a suffix to the name:


```
automata {
  Simple* = (takeTea->STOP).
  Two = (teaButton->Simple).
}
```

The informal meaning of events will in part be formalised by our definition (to be given later) of parallel composition. Informally we need to think of our events as *hand-shake events*, i.e. event that can be blocked or enabled by the context in which they execute. For example the **teaButton** event of a vending machine can only occur when some agent actually pushes the button. The pushing of the button can also be modelled by the **teaButton** event.

4.2 Event choice

A vending machine that has two buttons one for coffee the other for tea offers the user the *choice* to push either button. We model this using the *choice operator* $_|_$.

```
CM = (teaButton->takeTea->STOP|coffeeButton->takeCoffee->STOP).
```



this automata *branches* at the initial node.

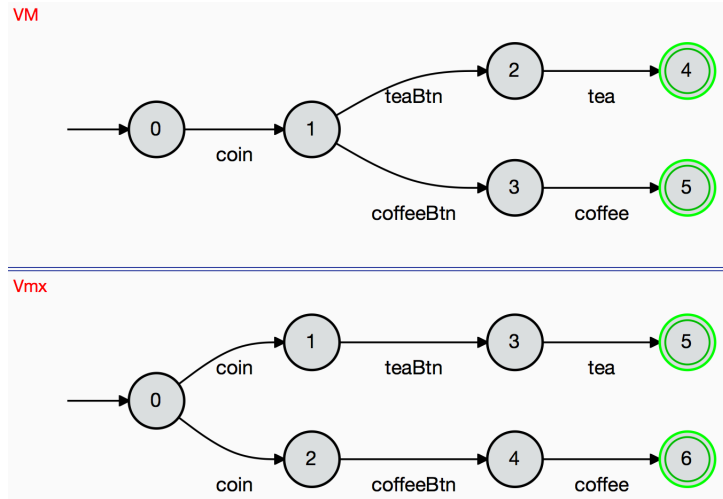
4.2.1 Non deterministic processes

The two processes VM and VMx both represent a vending machine that offers two drinks, **tea** and **coffee** after a coin is inserted. The two terms are different and they are represented by different automata.

```

automata
{
  VM = coin->((teaBtn->tea-> STOP)|(coffeeBtn->coffee->STOP)).
  VMx = (coin->teaBtn->tea-> STOP)|(coin->coffeeBtn->coffee->STOP).
}

```



Are VM and VMx equivalent processes? Before you can answer this you must decide what it means for two processes to be equivalent and there is many reasonable answers to this. If we assume either that the processes generate events or that they are used to recognise a sequence of events then the processes can reasonably be viewed as equivalent as both generate (recognise) the same two event sequences:

1. coin,teaBtn,tea
2. coin,coffeeBtn,coffee

But what if you were interacting with these processes and you wanted **coffee** then with the first machine you could always insert a **coin** than push the **coffeeBtn** and you would be able to get your **coffee**. In contrast with the second machine after inserting the **coin** you would not be able to push the **coffeeBtn**. Hence you would be able to distinguish the two processes. It is this notion of indistinguishable by any test that we are interested in here.

4.2.2 Properties of Choice

Choice is symmetric in that $X|Y$ is equal to $Y|X$. In what follows we write process equality as \sim and hence:

$$X|Y \sim Y|X \quad X|(Y|Z) \sim (X|Y)|Z$$

details about, \sim , process equality follow much later.

The Choice between two identical processes is no real choice and hence:

$$X|X \sim X \quad X|\text{STOP} \sim X$$

Meaning has been given to Processes in the form of an Axiomatic Algebraic semantics see [BW90]. Here we give meaning to processes by defining their operational behaviour as automata. From the details of what constitutes an automata we can see that the above equalities hold for our semantics. Looking at the same thing the other way in [BW90] they first define the algebraic properties of processes and then show the automata satisfy them.

4.3 Internal Choice +

Defined in CSP and ATP but not in ACP nor CCS. Used to represent a process that internally, with no external involvement, can choose to behave like either process. Consequently such a process can not be guaranteed to behave like either of the processes.

The relation between the two definitions of choice is captured by the axioms:

$$(a- > A + a- > B) =_F (a- > A|a- > B) =_F (a- > (A + B))$$

NOT bisim need to use testing or Failure equality!

Internal choice is best modelled by allowing automata to have a set of start states. Using a set of start states facilitates the modelling of processes with either probabilistic choice or *active* actions.

4.3.1 Properties of Internal Choice

Choice is symmetric in that $X+Y$ is equal to $Y+X$. In what follows we write process equality as \sim and hence:

$$X + Y \sim Y + X \quad X + (Y + Z) \sim (X + Y) + Z$$

4.4 Sequential composition

Two processes can be composed in sequence using \Rightarrow and sequential composition satisfies:

$$((a- > \text{STOP}) \Rightarrow A) \sim (a- > A), \quad X \Rightarrow (Y \Rightarrow Z) \sim (X \Rightarrow Y) \Rightarrow Z$$

$$(\text{STOP} \Rightarrow A) \sim A, \quad (A \Rightarrow \text{STOP}) \sim A$$

$$(X|Y) \Rightarrow Z \sim (X \Rightarrow Z)|(Y \Rightarrow Z), \quad (X + Y) \Rightarrow Z \sim (X \Rightarrow Z) + (Y \Rightarrow Z)$$

$$W \Rightarrow (X + Y) \sim (W \Rightarrow X) + (W \Rightarrow Y) \sim (W \Rightarrow X)|(W \Rightarrow Y)$$

4.5 Non terminating processes


We call the set of know processes the *Process name space*. Initially the *Process name space* is $\{\text{STOP}\}$. Each process definition $P1 = \dots$ adds the the defined process $P1$ to the name space.

Processes consist of a set of states, an initial state and a set of event labelled state transitions. Given a process has a set of states and a set of transitions it is reasonable that the process can be *conceptual identified* with its initial state.


Clearly any state S in process $P1$ could also be *conceptual identified* with the the process consisting of the same set of states and transitions but with initial state S . We use this idea to define non terminating processes. By allowing any valid process to be used where $\{\text{STOP}\}$ has been used we can define non terminating or cyclic processes.

To build events that do not terminate we can replace STOP with the name of the process we are defining thus $T = (\text{takeTea} \rightarrow \text{STOP})$. becomes $Tt = (\text{takeTea} \rightarrow Tt)$. The process Tt can endlessly perform the takeTea event.

$Tt = (\text{takeTea} \rightarrow Tt)$.




$BT = (\text{teaButton} \rightarrow \text{takeTea} \rightarrow BT)$.



We allow *local process* or states to be defined within a process definition by separating definitions with a comma. The local process do not appear in the *Process name space*.

$P = (a \rightarrow Q)$,
 $Q = (b \rightarrow P \mid c \rightarrow Q)$.




This allows complex processes to be defined without cluttering the *Process name space*.

4.6 Translating any finite state automata into a process term

It is often easy to sketch your understanding of a processes behaviour as an automata. Then from any automata we can construct the process term with the behaviour given by the automata. Our tool will automatically generate the automata from the term. The generation of the term from the automata can be achieved quite mechanically as follows:

1. name all nodes (or all nodes with more than one in and one out event) with a process name
2. define each of the processes and the choice of events leaving them
3. end each process definition with a comma except for the last process that must end with a full stop.



For the above automata node 0 we name **TrRed** and node 1 we name **TrGreen**. Then we define the events leaving these nodes

$$\begin{aligned} \text{TrRed} &= (\text{red} \rightarrow \text{TrRed} \mid \text{turnGreen} \rightarrow \text{TrGreen}), \\ \text{TrGreen} &= (\text{green} \rightarrow \text{TrGreen} \mid \text{turnRed} \rightarrow \text{TrRed}). \end{aligned}$$

The result of this construction is the definition of the first process **TrRed**, all other processes, in this case just **TrGreen**, are *local* definitions.

5 Semantics of processes with atomic events

Distinct automata can be used to represent exactly the same process. We formalise this by defining a semantic equivalence of automata. The question as to what automata should be equated and how do you justify your notion of equality we leave to later. Here we are going to introduce two very different notions of process equality that are widely used and widely discussed in the literature.

5.1 Complete Trace equality

The trace semantics of a process are the set of executions the process can undertake.

Complete finite traces must end in state from which no event can occur:

$$\text{Tr}_{Fin}(P) \triangleq \{tr : \exists n : S_P \xrightarrow{tr} n \wedge \pi(n) = \emptyset\}$$

Infinite traces do not end:

$$\text{Tr}_{Inf}(P) \triangleq \{tr : S_P \xrightarrow{tr} \}$$

The complete traces of a process

$$\text{Tr}_c \triangleq \text{Tr}_{Fin} \cup \text{Tr}_{inf}$$

Complete trace equality:

$$P =_{\text{Tr}_c} Q \triangleq \text{Tr}_c(P) = \text{Tr}_c(Q)$$

Trace equality does not distinguish deterministic processes from non deterministic processes and hence the two processes in Section 4.2.1 are identified.

5.2 Bisimulation

A bisimulation \sim is relation on the nodes of an automata that is symmetric $n \sim m \Rightarrow m \sim n$ and

$$n \sim m \wedge n \xrightarrow{a} n' \Rightarrow \exists m'. m \xrightarrow{a} m' \wedge n' \sim m'$$

Two processes P and Q are bisimilar if and only if there is a bisimulation relation that relate their start nodes, $P_S \sim Q_S$.

Bisimulation equivalence is much finer than complete trace equality and only equate processes that could not possibly be distinguished. Consequently the two processes in Section 4.2.1 are not bisimilar (bisimulation equivalent).

To help us understand how bisimulation equivalence works we give a simple co-inductive algorithm to compute the maximal bisimulation relation using a node colouring where nodes with the same colour are related.

1. Initially colour all nodes with the same colour.
2. Repeatedly recolour the nodes using

$$Col_{i+1}(n) \triangleq \{(a, Col_i(m)).n \xrightarrow{a} m\}$$

3. stop when the recolouring changes nothing.

$$Col_{i+1}(n) = Col_{i+1}(m) \Leftrightarrow Col_i(n) = Col_i(m)$$

$$\text{Colour map } Col(n) \mapsto \{(a, Col(m)).n \xrightarrow{a} m\}$$

The maximal bisimulation relation that can be computed very quickly and bisimilar nodes, nodes with the same colour, can be identified to produce a simpler automata. Using this colouring algorithm it is easy to see the result of applying bisimulation simplification to simple automata.

This algorithm can be applied to many automata at the same time and can be used to compute an equivalence class on a set of automata. To cope with the multiple start states we first compute the bisimulation colouring and Colour Map. Then two processes, A and B are bisimilar if and only if

$$A \sim B \triangleq \bigcup_{s \in A_S} Col(s) = \bigcup_{s \in B_S} Col(s)$$

Bisimulation has attractive mathematical properties, is easy to compute and has proven to be of practical use. Consequently bisimulation relations have been defined on many different structures.

Both these equivalences are congruent with respect to our process operators (substitution of equivalent processes).

$$A \sim B \Rightarrow A || P \sim A || P$$

$$A =_{Tr_c} B \Rightarrow A || P =_{Tr_c} A || P$$

5.3 Failure Semantics

This denotational semantics is taken from [?] where it has been applied to processes with hand shake events. It has been shown to correspond to a natural notion of Testing semantics in [].

$$Fail(A) \triangleq \{(t, R).s \xrightarrow{t} A \wedge R \subseteq (Act - \pi(n))\}$$

$$A =_F B \triangleq Fail(A) = Fail(B)$$

5.4 Quiescent Trace semantics

This denotational semantics is taken from [?] where it has been applied to processes with broadcast events. It has been shown to correspond to a natural notion of Testing semantics in [?].

$$\text{Qtr}(A) \triangleq \{t.s \xrightarrow{t}_A n \wedge \pi(n) \cap \text{Act}_! = \emptyset\}$$

$$A =_{\text{Qtr}} B \triangleq \text{Qtr}(A) = \text{Qtr}(B)$$

Note $A =_F B \Rightarrow A =_{\text{Qtr}} B$


6 Event hiding and process simplification

Our process language can be used both to specify *implementations* of processes and more abstract *specifications*. Our tool is designed to support the approach popularised by Event B. Both the abstract specification and the more concrete implementation are specified in the same language. The tool then checks that the implementation is a *refinement* of the specification.

In our event based approach events in the implementation but not the specification are hidden by abstraction and the resultant process is checked for equality with the specification. In this section we define event hiding and abstraction.

We can make events private by hiding them so they can not be seen. $_{\setminus\{t\}}$ operator renames the t event to the unobservable **tau** event.

See following example:



In the above example **abs** introduces two observable events:

1. $0 \xrightarrow{a} 2$ in place of the event sequence $0 \xrightarrow{a} 1$ and $1 \xrightarrow{\tau} 2$
2. $1 \xrightarrow{b} 4$ in place of the event sequence $1 \xrightarrow{\tau} 2$ and $2 \xrightarrow{b} 4$

The **abs**($_{\setminus}$) operator abstract away the **tau** events. To cope with cases where many τ events can be executed one after the other we first define $x \xrightarrow{\tau} y$.

A sequence of zero or more τ events

$$x \xrightarrow{\tau} y \triangleq \exists i \geq 0 : \exists n_1, n_2, \dots, n_i : x \xrightarrow{\tau} n_1, n_1 \xrightarrow{\tau} n_2 \dots n_i \xrightarrow{\tau} y$$

can be executed unseen and are represented as $x \xRightarrow{\tau} y$. When $i = 0$ we have $x \xRightarrow{\tau} x$ for any x .

Abstraction constructs, $x \xRightarrow{a} y$ the observable semantics:

$$x \xRightarrow{a} y \triangleq \exists u, v : x \xRightarrow{\tau} u \wedge u \xrightarrow{a} v \wedge v \xRightarrow{\tau} y$$

6.1 Event hiding and non terminating processes

The literature is divided on how to hide τ events that loop. CSP refers to these processes with τ loops as *diverging* and models them as having potentially *chaotic* behaviour. CCS and Discrete Event Systems DES, assumes them to be benign as simply prunes them. Here we offer both options. The CSP option assumes that the system can behave *unfairly* and the CCS option assumes the system behaves *fairly*.

The command `abs(_)` is based on the fair assumption and `abs{unfair}(_)` is based on the unfair assumption.

With the unfair assumption congruence with respect to interleaving parallel composition requires some care because the relation between events from parallel components is *fair*. That is to say the infinite execution of an event from one process can not prevent the parallel process executing an unrelated event.

In other words one process, P, diverging will not effect the events of a second process, Q, running in parallel whereas with the unfair assumption it can block other events on P. But with interleaving composition $P \parallel Q$ the events of P and Q can no longer be distinguished.


To accommodate this label nodes on a τ loop as divergent and require that parallel composition preserves divergence. This has the effect of allowing divergence of P to block the events of Q.

bug to be checked

```

automata {
  S* = a->X,
    X = (t->X|b->STOP).
  Simple = S\{t}.
  O* = Simple||(d->STOP).
  One = abs{unfair}(O).
  T* = abs{unfair}(Simple).
  Two = T||(d->STOP).
}

```



Pragmatically divergence and deadlock are rarely wanted and their existence merely indicate that the definitions are erroneous. In such situations it is unimportant how we model divergence as it will be removed. In situations where we do want to model process with deadlock or divergent behaviour then we need to be more careful.

Frequently we are interested in how a system behaves when errors occur but when errors occur is rarely determined exactly. Hence modelling a systems correct and error behaviour will introduce some probabilistic or non deterministic behaviour. Hiding both errors and their handling may introduce divergence and this may be best interpreted *fairly*.

7 Concurrent Processes

So far we have have defined sequential processes but now we wish to define how two sequential processes behave when they are both run together. This is modelled using the parallel composition operator $- \parallel -$.



Two processes run in parallel can only interact via event synchronisation. How this is defined depends upon the interpretation you wish to give your events. In what follows we consider two distinct styles of events, hand shake events as found in CSP and CCS as well as broadcast events as found in IOA.

Two processes run in parallel can only interact via event synchronisation and events on only synchronise with other event having the same name.

Below left the parallel composition of two simple processes $(a \rightarrow \text{STOP} \parallel b \rightarrow \text{STOP})$ is shown to be semantically equivalent to a single sequential process $(a \rightarrow b \rightarrow \text{STOP} \mid b \rightarrow a \rightarrow \text{STOP})$. The ability to replace parallel processes with sequential process greatly simplifies their analysis.

$$P = ((a \rightarrow b \rightarrow c \rightarrow \text{STOP}) \parallel (x \rightarrow y \rightarrow z \rightarrow \text{STOP})).$$

$$(a \rightarrow \text{STOP} \parallel b \rightarrow \text{STOP}) \sim (a \rightarrow b \rightarrow \text{STOP} \mid b \rightarrow a \rightarrow \text{STOP})$$



Above right we have two processes each with three events and no two event have the same name hence the event from each process can be **interleaved** in any way.


This definition of parallel composition equates parallel processes with a single sequential process. As such some aspects of real parallel processes are lost. Some times this equality is referred to as being based upon the *interleaving assumption*.

Without synchronization two processes are independent and hence their events interleave and the state space of the composition of the processes is the product of the state space of the constituent processes.

7.1 Handshake synchronisation

In the Process tool events from different concurrent processes that have the same name must synchronize and only these events synchronize. That is neither process can execute the synchronising event on its own. These synchronising events are only executed when both processes are ready to execute them. Below only differs from the previous process in that the second event in both processes has the same name and hence must synchronize and the resulting m event is then hidden (renamed τ).

$$P = ((a \rightarrow m \rightarrow c \rightarrow \text{STOP}) \parallel (x \rightarrow m \rightarrow z \rightarrow \text{STOP})) \setminus \{m\}.$$



Event synchronization is the only mechanism for concurrent process interaction and because of event synchronisation we know:

If you can see and event you can synchronize with it and you can block it.

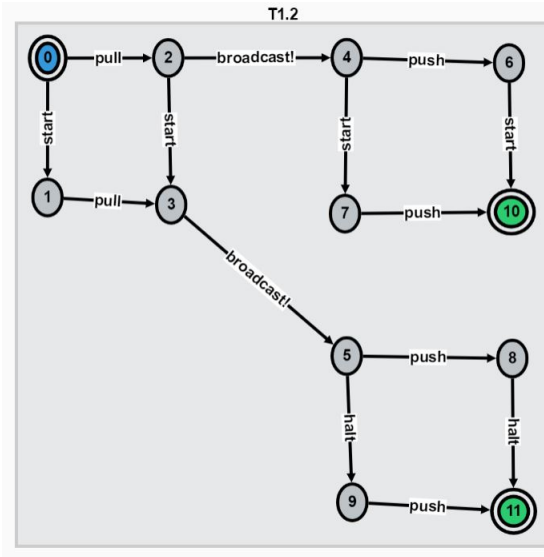
Hence the only way the control the order of two events from different concurrent processes is to introduce a synchronizing event. In above the **a** event and the **z** event are from different concurrent processes in the interleaving example either could occur first. Whereas in the synchronization of the **m** events forces the **a** event to occur before the **z**.

Another effect of synchronization is to reduce the size of the reachable state space of the automata. Note the first two events **a** and **x** can be performed in either order but only when both **a** and **x** have been performed and both processes are ready to perform **b** dose the **b** event actually get performed.

7.2 Broadcast event synchronisation

In the previous section when events from two processes synchronised the the synchronising events from both processes were treated the same, both could block the other process. This style of synchronisation captures real events such as **pushing a button** I cannot push a button that is not there or is frozen nor can button on a vending machine be pushed if I am not prepared to push it. Both me and the vending machine must wait for the other to be ready before the **buttonPush** event can occur.

$(a \rightarrow \text{broadcast!} \rightarrow c \rightarrow \text{STOP}) \parallel (x \rightarrow \text{broadcast?} \rightarrow y \rightarrow \text{STOP})$



send radio warning	warning!	green light is shining	green!
hear radio warning	warning?	I see the green light	green?

point to point Emails are often messages from one person to one other unique person.

Note: if you are not listening you might miss the warning and if you are not looking you might not see the green light. Listening/looking is formalised as being in a state from which the receive event is enabled. We model processes so that if you are listening/looking then you will hear/see the broadcast event.

Consequently wo **broadcast?** events will not synchronise like hand shake events but behave as shown below:

The automata can be seen below:




8 Renaming events and simplifying Processes

Frequently when we are thinking of one process it is natural to give an event a particular name. But, when considered from the perspective of another process that may interact with it this name might be confusing. Consequently we introduce ways to rename events. Finally we show how bisimulation colouring can be used to simplify processes. To aid understanding we will consider a simple buffer example in the following.

8.1 Labelling Processes


In the following example we make use of a one place buffer `Buf` is a process that when empty can receive some thing `in` and when full can return it `out`.



By labelling processes `one:Buf` the tool labels all events in the process `one.in` and `one.out`.

Using process labelling we can make two differently label copies of a process and compose them in parallel to build the interleaving of the two copies.

$$B2 = (\text{one:Buf} \mid \mid \text{two:Buf}) .$$




8.2 Event renaming

If two events from processes run in parallel have the same name they, and only they, must synchronise.

Pragmatically when you compose two processes in parallel you should check the name of events you want to synchronise and where necessary rename them to enforce the desired synchronisation.

We force the synchronisation of the output from buffer `one` with the input to buffer `two` by event renaming.


$$B3 = (\text{one:Buf} / \{\text{move}/\text{one.out}\} \mid \mid \text{two:Buf} / \{\text{move}/\text{two.in}\}) .$$




Note that the result is much simpler than the interleaving as the `move` event now can only occur when **both** buffers are able to perform it.

8.3 Process simplification

We can go further and hide the `move` event by applying $_ \backslash \{\text{move}\}$. The `move` event becomes a `tau` event that can neither be synchronized with nor blocked.



The `tau` events can be removed by **abstraction**, (the application of $\text{abs}(_)$) otherwise known as building the *observational* semantics. With a little effort nodes, 1 and 2 in B5 can be seen to be essentially the same. They are actually bisimilar but we will not be going into details here. These nodes can be identified to produce a simpler but equivalent automata by the application of $\text{simp}(_)$.



Event hiding is commonly, but not exclusively, used to model private communication.

9 Testing and equality

In this section we are going to introduce a very generic notion of equality, testing equality. This can be applied to to any set of things all they need is an operator to compose them and a definition of how to observe them, here the set of things is our set of processes, the operator to compose then is parallel composition and finally when we observe a process all we can see is the complete trace of events that are executed.

When a non-deterministic system is tested it is placed in some context and the combination of the system and text context is executed, but this test must be run a number of times and the set of observations (results) recorded.

Our processes, E , are taken from a set of processes \mathbb{E} . A context, $_ \parallel X$, consists of a process, $X \in \mathbb{E}$, run in parallel with the process under test. Let $\Xi \triangleq \{ _ \parallel X \mid X \in \mathbb{E} \}$ be the set of all contexts. Placing E in context X can be written as $[E]_X$ or as $X \parallel E$. A single experiment consists of observing a single execution of $[E]_X$ and results in a single trace, taken from a set of possible observations Tr^c , being recorded. For non deterministic processes the experiment must be repeated and a set of observations Tr^c , being recorded.

A specification can be interpreted as a *contract* consisting of the *assumption* that the process will be placed only in one of the specified contexts Ξ and a *guarantee* that the observation of its behaviour will be one of the observations defined by the mapping $O : \mathbb{E} \rightarrow \Xi \rightarrow \wp Tr^c$. The mapping O defines what can be observed for all processes in any of the assumed contexts. Hence for any fixed Ξ we have a definition of the semantic equivalence of processes.

Definition

$$\llbracket A \rrbracket_{\Xi, O} \triangleq \{ (x, o) \mid x \in \Xi \wedge o \in O([A]_x) \}$$

and equality is

$$A =_{\exists, O} C \triangleq \llbracket C \rrbracket_{\exists, O} = \llbracket A \rrbracket_{\exists, O}$$

Let \mathbb{E} be the set of LTS and Tr^c be the complete traces of an automata then all we need to define to fix our definition of testing equality is the definition of \parallel parallel composition. \square

We have defined parallel composition between handshake events and another definition of parallel composition for broadcast events. Using parallel composition for **handshake events** the above definition of testing equality has been shown to be the same as the well known **Failure equality** from CSP. Whereas using parallel composition between **broadcast events** the above definition of testing equality corresponds to the well known **quiescent trace equality**.

10 Indexed Process definitions

Basic process definitions you have seen so far a fixed finite set of states. This accurately reflects many situations very well and allows easy and complete push button verification. Alternatively when what you are modelling has infinite state or an unknown state size you could use symbolic models but verification frequently requires input from a domain expert and is very time consuming.

The approach adopted here is to define both states and events using an index and limit the size of the index by a parameter. Prior to using a parameter it must be declared and given a fixed value. A finite state approximation of indexed process can be built and size of the approximation can be changed by changing the declared value of the parameter.

10.1 The small world assumption


Most program bugs can be found while restricting variables to range over a small domain. Using this assumption we model processes with variables by indexing the processes and restricting the indexes to range over a small domain. Having done this the variables in the state can be removed by instantiating the variables with values from the small domain.

More than one parameter can be used to define a process. Once all the parameters are fixed you are back to a basic process with a finite set of states and events. Processes can be indexed in different ways to achieve conceptually different things. The first we consider is how to build a process of parametrised size, the second is to model events that input or output data and finally how to model a parametrised number of concurrent processes.

10.2 State indexing

We can define a process consisting of an unknown number of states. To do this we must index the local states (or local processes). We will consider a simple **Purse** that can contain a number of **coins**. We define the **Purse** based on a parameter **N** that depicts its size.

Automata for a **Purse** that can contain 3 coins.



The first thing we do is define a constant to be used for the size of the automata to be constructed:

```
const N = 3
```

Next we define the automata:

```
automata {
  Purse = P[0],
  P[c:0..N] = (when c<N  coin -> P[c+1] |
               when c>0  spend -> P[c-1]).
}
```

The first line `Purse = P[0]`, defines that the purse is initially empty then the definition `P[c:1..N] =` defines the N processes `P[1]`, `P[2]` and `P[3]`

The term `P[c:1..N]` on the left of the equality can be thought of as assigning a value to a variable `c`. The term `P[c+1]` on the right of the equality reads the value in the variable and then "returns" to the purse in a new state where "`c:=c+1`".


On the right hand side of the equality we define guarded events:

```
when(c<N) coin->C[c+1]
```

the `coin` event will only occur when the guard is true, `c<N` and the event ends at node `C[c+1]`. **Note a guard only applies to one event.** Each time you add a choice you need to add any required guard.

The command `when(c<N) P[0]...` uses process `P[0]` that is not event prefixed. Hence will this command not compile. Whereas `if...then...else P[0]` will compile and produces the automata displayed.

```
PurseIf = P[0],
P[c:0..N] =
  (if (c < N) then  addcoin-> P[c+1]
   else    P[0] ).
```




In Section 10.3 we will make use of this property of `if` to simplify the definition of process. An interesting exercise is to try to define the process with out using the `if` command.

10.3 Event indexing

An indexed event can be used to model events that input or output values.

A one place buffer that can accept as input a number from the range $1..N$ and then must output that value is can be represented by an automata with $N+1$ states.



The buffer is defined by:

$\text{Buf} = \text{input}[v:1..N] \rightarrow \text{out}[v] \rightarrow \text{Buf}.$

The event $\text{input}[v:1..N]$ **declares** a new variable v and when it executes it **inputs** a value that is assigned to the variable. The subsequent event $\text{out}[v]$ refers to the previously declared variable v and when it is executed it **outputs** the value held in it. So information flows from the **input** event to the **out** event and hence the names.

10.4 Cars Example


This example requires both state and event indexing.

A car can travel at different speeds and can accelerate at different rates. But no matter how hard it tries to accelerate it can never go beyond its maximum speed.

In this example the car is indexed by its **speed** and the number on the nodes corresponds to the speed of the car. The rate of acceleration is indexed by **a** and this index is declared in the event **accelerate**. Hence the rate of acceleration can not appear as an index to the car as that is defined prior to the definition of the event.

```

const N = 3
automata }
Car = C[0],
  C[speed:0..N] =
    (when speed < N accelerate[a:1..N-1] ->
      ( if (speed+a < N) then C[speed+a]
        else C[N])
    | when speed > 0 brake -> C[speed-1]).
}
  
```



In the example above the Car has a maximum speed of 3 and a maximum acceleration of 2. But when the Car has speed of 2 the effect of accelerating at 2 is only to change the speed by 1.

10.5 From Natural Language to indexed process model

Natural languages are expressive but ambiguous. Added to which we are interested in describing event based models and there is no one universal way to describe such systems. This leads to many problems, some of which can be overcome by breaking the task of formalising these informal specifications into some simple steps.

Step 1 find indexes and indexed states

Step 2 find indexed events

Step 3 find all events

Step 4 build automata either sketch and code or code and view.

Step 5 inspect the automata and validate it against specification

We will demonstrate this with a simple example of a lockable door.

Closed doors are always locked. The door starts closed. The lock can hold any of a number of codes. To open you need to input the correct code and after opening the door can only close. Inputting the wrong code is an error and the door returns its start state. Before using the door the code must be set.


You may assume that only an administrator can set the code where as any one may use the door by entering codes but such distinctions are not part of the model.

Step 1 When the numbers of states or events is not fixed but is dependent upon some parameter then you need to build what we call an indexed process. The parameter is an index and in our example this is the `code` the lock uses. As the code needs to be stored by the Lock we need indexed states `L[j:1..N]`.

Step 2 There are two indexed events `setlock[k:1..N]` to set the state of the Lock and `enter[ji:1..N]` to enter a code when trying to open the door.

Step 3 The list of all events: `open,close,error,enter[]` and `setlock[]`

Step 4 Automata when `N==2`



This is defined by:

```
Lock = (setlock[k:1..N] -> L[k]),
L[j:1..N] = (enter[i:1..N] ->
    ( when (i==j) open ->close->L[j]
      | when(i!=j) error->Lock)).
```


Step 5 Note the value input in the `setlock[i:1..N]` event is stored in the state of the process `L[i]` for subsequent comparison with the value input in the `enter[i:1..N]` event.


10.6 Indexing concurrent processes.

This means producing an indexed number of concurrent processes but has been implemented in a very restricted. As implemented the indexed process can only communicate with processes that are running in parallel with them. Alas they cannot communicate with each other.

If you want N **Worker** processes, each labeled with `[1]`, `[2]`, \dots `[N]`

```
Worker = (getTask -> doTask -> Worker).
Workers = (forall [i:1..N] ([i]:Worker)).
```






We can add a **Farmer** process to hand out the **Tasks** to the **Workers** in order. Then build a **Farm** composed of the **Farmer** and the **Workers**.

```
Farmer = F[1],
F[i:1..N] = (when (i<N) [i].getTask->F[i+1]
            | when (i>=N) [i].getTask->F[1]).
```

```
Farm = (Farmer || Workers).
```



The **Farmer** process is far from ideal in some regards.

11 Petri Nets

What we refer to as a Petri Net is actually an extension of Petri Nets with labels on the transition. This appear in the literature under the name of the *Box calculus*. Net semantics differ from automata significantly in that they offer a **non interleaving** interpretation of parallel processes. Such semantics are frequently referred to as a **true concurrency semantics**. We adopt a definition of parallel composition that originates in the Box calculus and is not a part of main stream Petri Nets.

A Petri Net is a tuple $A \triangleq (P_A, S_A, T_A, \text{Arcs}_A, \alpha_A, \text{Name}_A)$ where:

P_A is the set of places,

S_A is the set of initial Markings $M_A \subseteq S_A$ where a Marking is a set of places $M_A \subseteq P_A$

T_A is the set of transitions

Edges_A Edges join node to transitions and transitions to nodes. There can be multiple places joined to a single transition and multiple transition joined to a single place.

$$\text{Edges}_A \subseteq \{(p, t) : p \in P_A, t \in T_A\} \cup \{(t, p) : p \in P_A, t \in T_A\}$$

α_A is the alphabet of the process

Name_A a transition naming function $\text{Name}_A : T_A \rightarrow \alpha_A$

□

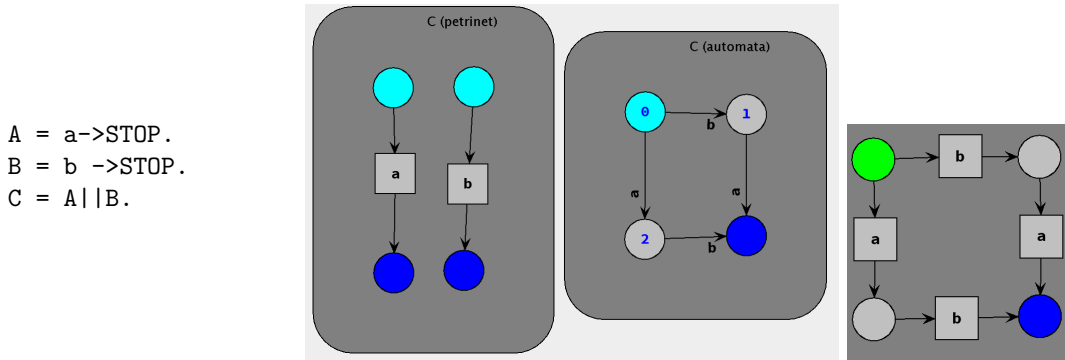
11.1 Appearance

Nets have Places, large circles, not nodes and the initial state of a Petri Net is a set of places each *marked* with a token, a small black circle. The transitions are represented by Boxes and event names. Arcs are added from *pre place*, $\bullet t$ to the Box of transition t and from transition t Box to *post place* $t \bullet$.

Finite state Petri Nets like finite state automata can be approximation of potential infinite state processes. The location of the Petri net transitions is the same as that for Automata.

There is a simple relation between finite state sequential automata and finite state Petri Nets. In addition the definition of event hiding and event renaming on Petri Nets and its relation to event hiding and event renaming on automata is quite obvious.

Our tool takes process specifications \mathcal{P} and builds finite state automata P . But in addition it can build Petri Nets from specification $\text{Petri}(\mathcal{P})$. PertiNets provide a richer process model than automata in that automata.



Two simple processes run in parallel are represented by a single process composed of the interleaving of the events. Whereas a Petri Net can be used to capture the concurrency with a net that has two sf tokens or to model the interleaving with a Net having a single token.

The operations defined on atomic processes can be lifted to operations on Petri Nets.

Parallel merge of two processes, parallel composition with no synchronisation is just the union of the component nets. Hence automata node is represented by a marking. That is a pair of places, one element of the pair taken from each component process. Thus each place represents the state of one of the component processes. Each process can be give a unique location and then each node annotated with the location of its process. The start node corresponds to the the pair of initially marked places.

Event synchronisation of t_1 , $\bullet t_1 \xrightarrow{a} t_1 \bullet$ and t_2 , $\bullet t_2 \xrightarrow{a} t_2 \bullet$ is a new transition t_3 with the same name and with pre places the union of the component transitions pre places. The post places are constructed similarly.

$$t_3 \triangleq (\bullet t_1 \cup \bullet t_2) \xrightarrow{a} (t_1 \bullet \cup t_2 \bullet)$$

11.2 Token Rule

We define **TokenRule** that maps Petri Nets to automata. A transition can only be executed when all its pre places are marked and the result of firing a transition is to move the tokens from the pre places to the post places.

This can easily be extended to allow a set of transitions to be executed at the same time when the multi-set union of the places are marked. Thus Petri Nets offer a **true concurrent** semantics of processes.

These are the same automata that would have been constructed had we built the automata directly from our process language, hence we have:

$$\text{TokenRule}(\text{Petri}(\mathcal{A})) \sim \mathcal{A}$$


Let us write A and B for Petri Nets. We have lifted operations Op_a on automata to operations Op_n on Petri Nets so that they obey the following algebraic rules:

$$\text{TokenRule}(A \parallel_n B) \sim (\text{TokenRule}(A) \parallel \text{TokenRule}(B))$$

$$\text{TokenRule}(A\$\{x\}) \sim (\text{TokenRule}(A))\$\{x\} \quad \text{TokenRule}(\text{abs}(A)) \sim \text{abs}_n(\text{TokenRule}(A))$$

As defined Petri Nets are semantically richer than automata as they tell us more about processes than automata can. But they do not capture everything from our process language see below:

$P = (a \rightarrow x \rightarrow b \rightarrow \text{STOP} \parallel$
 $\quad r \rightarrow x \rightarrow t \rightarrow \text{STOP}).$
 $Q = (a \rightarrow x \rightarrow t \rightarrow \text{STOP} \parallel$
 $\quad r \rightarrow x \rightarrow b \rightarrow \text{STOP}).$



In the example above the processes P and Q are not the same as the events b and t *belong* to different processes and this information is lost in the Petri Net semantics.

We can extend both automata and Petri Nets to record the location, or ownership, of events. This way we have located automata are equivalent to located Petri Nets.

11.3 From Automata to Petri Nets

Building Petri Nets from automata with located transitions **Owners Rule** works as follows:

Automata alone do not preserve concurrency information. By recording

Define a Location space where each location is a subset of primitive locations. Build a Place for each primitive location *own* by building a projection of the automata onto the location. In essence events at a location containing *own* are preserved all other events are hidden. Finally the set of projected automata can be composed in parallel synchronising only the events that have been projected for the same underlying event.

The algorithm we use is:

```

for each primitive location 0
push rootnode onto stack
while stack not empty
  pop node from stack
  if not processed
    collect set S of all nodes connected (undirected) by edges not at 0
    build Place and mapping from all nodes in S to Place
    for all nodes connected by edges at 0 push toNode onto Stack
push rootnode onto stack
while stack not empty
  pop node from stack
  if not processed
    for all nodes connected by edges at 0 build transition between mapped Places

```

The round trip from an automata to a Petri Net (Owner Rule) and back (Token Rule) produces a bisimilarly equivalent automata. This has been tested for thousands of automata.

Two semantically equivalent ways to view the same process

The advantage of being able to model the same process in two quite distinct styles allows us to define operations on one and project them to the other.

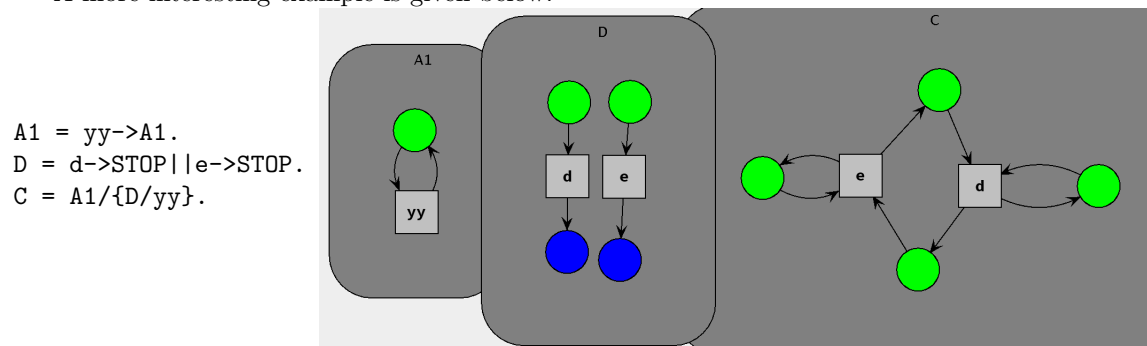
Process equivalences, bisimulation, failures equality and complete trace equality can all be defined on automata and easily extended to our **located automata** by the simple trick of viewing the location *s* a part of the event name.

Whereas event refinement can be defined on Petri Nets. Because of the addition of location the two models are equivalent but nonetheless trying to define event refinement on located automata or testing equivalence on Petri Nets is far from easy. Several definitions of process equality have been defined on Petri Net semantics but they are difficult to understand and very hard to compare thoroughly.

Event Refinement

We can define how to refine a single event in a process (PetriNet) into a whole process (Petri Net).

A more interesting example is given below.



To check your intuitions about Petri Nets ask yourself if the two "*wings*" places which are always marked are just an artefact of our algorithm and are semantically redundant or not?

11.4 Advantages of mixed semantics

Petri Nets can be very much smaller than automata and also support event refinement. Automata support a rich set of equivalences that can be justified by plausible testing semantics and easily represent reachability of events.

Two options for the Hiding and simplifying Nets

1. define abstraction and bisimulation on nets OR
2. Use the **Token Rule** to map a Petri Net to an Automata, simplify the automata and then use **Owners Rule** to map the result back to Petri Nets.

12 Semantic equivalence of processes

One of the key questions is: what automata should be equated and how do you justify your notion of equality? Before we try to answer this we need to remember what equality means.

12.1 Classic notions of equality

An equality is a relation between items from some domain, here we are interested in process equality so the domain is the set of all processes. Equality has the following properties:

1. an equivalence relation, a reflexive, symmetric and transitive relation
2. substitutive and
3. where equivalent objects are indistinguishable.

Indistinguishable processes can be formalised as *testing equivalent*. Substitutive can be formalised as congruent w.r.t. process operators, from a practical perspective the most important being parallel composition.

12.2 Complete Trace equality $=_{Tr^c}$ or in the tool $_ \# _$

See Section 5.1 for details.

This can be computed by the well documented *Subset Construction* that computes a Deterministic automata from a non deterministic automata. and then applying the bisimulation algorithm discussed in Section 5.2

12.3 Failure equality $=_F$ or in the tool $_ * _$

The Failure semantics of a process consists of a set of Failures. A Failure is a pair consisting of a trace and a Refusal set (tr, R) .

$$Fail(P) = \{(\rho, R) : S_P \xrightarrow{\rho} n \wedge R \subseteq \overline{\pi(n)}\}$$

If you execute a process in parallel with a context then the context can:

1. see the trace of execution and

2. when the execution stops it can infer that the process refuses to perform any of the events that the context is offering.

Processes are *Failure* equivalent if they have the same set of Failures.

$$P =_F Q = \text{Fail}(P) = \text{Fail}(Q)$$

Failure refinement, \sqsubseteq_F

$$P \sqsubseteq_F Q = \text{Fail}(P) \supseteq \text{Fail}(Q)$$

One of the most important aspects of Failure semantics is that failure refinement is the same as testing refinement using parallel composition for handshake events, \sqsubseteq_{Ths} .

$$P \sqsubseteq_F Q = P \sqsubseteq_{Ths} Q$$

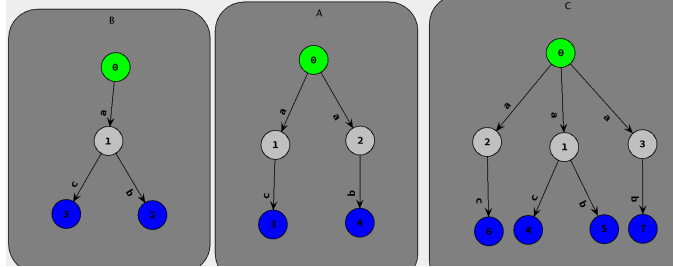
12.4 Bisimulation

See Section 5.2 for details.

12.5 Example processes

Process equalities can be checked by the tool by defining operations after the processes have been defined and the automata built. Not equivalence is defined by prefixing one of the equality operators with a !.

```
processes {
  A=(a->b->STOP|a->c->STOP).
  B=a->(b->STOP|c->STOP).
  C=(a->b->STOP|a->c->STOP|
    a->(b->STOP|c->STOP)).
}
automata A,B.
operation {
  A # B. A !* B.
  A * C. A !~ c.
}
```



The operations state that A and B are trace equivalent but not failure equivalent. Whereas A and C are Failure equivalent but not bisimilar equivalent.

A might perform an *a* event and then fail, when offered to synchronise with a *b* event. In contrast whenever B performs an *a* event it will never refuse to perform an *b* event. Consequently the two processes are distinguishable by testing.

Although A and C are not bisimilar equivalent they can not be distinguished by any test.

Some relations between equalities can be used to understand certain aspects of the equivalences. Bisimulation is strictly stronger than Failure equality which in turn is strictly stronger than Complete trace equality.

$$P \sim Q \Rightarrow P =_F Q \Rightarrow P =_{Trc} Q$$

A deterministic process P is failure equivalent to another process only when it is bisimilar to it.

$$det(P) \Rightarrow (P \sim Q \Leftrightarrow P =_F Q)$$

Consequently if one process is deterministic then failure equality, refinement can be computed easily by computing bisimulation.

Weak equalities, refinements can be computed by first computing the weak semantics and then, computing the strong equality or refinement on the weak semantics we have the weak semantics on the original automata. A useful analytic technique is to start with a detailed deterministic process, hide some events and simplify the process using bisimulation. This builds a more abstract representation that may well be deterministic, if it is not deterministic this tells you some thing about your detailed process.

13 TLA+ Overview

In what follows we will refer to a formal step that moves from an abstract specification to more concrete specification as a refinement step and the reverse step as an abstraction step.

TLA+ is mainly used to model shared memory concurrency but not exclusively. A central aspect of TLA+ specifications is use of logic to specify assignment and behaviour.

The basic building blocks of TLA+ are Modules. Modules contain two basic components variables and actions. The state of a module is an evaluation, a variable value mapping $State : Var \rightarrow Val$. That is the state is a context in which expressions can be evaluated. Commonly State is modelled as a record where the names of the fields in the record are the names of the variables. The state of a module can be changed by executing one of its named actions.

TLA+ has no parallel composition operator Modules are commonly defined to model both the process under construction and the world around it. In Process algebra an observable event needs to synchronise with its counterpart in another process, as such the two component events can be thought of as *half events*. This is not true for the actions of TLA+.

Assignment is defined as a predicate by introducing primed representation of pre state. This way $x := x+1$ becomes $x' = x+1$. Converting as much as you can into predicates makes the use of a theorem prover more effective.

Actions consist of a guard and a set of assignments. Consequently the semantics of an action can be given as a named relation over the lifted state, $State^\perp$. To facilitate reasoning actions are defined as a conjunction of predicates. Hence an action is a named predicate.

To define a TLA+ action that sends data we write $Send(d)$ and

$$Next \triangleq (\exists d \in Data : Send(d)) \wedge Rcv$$

Although this action resides solely on the sending process it may be in a Module that specifies both the sending and the receiving processes. The receiving action makes no explicit reference to the value d because TLA+ uses shared memory all that need happen is that

$$Send(d) \triangleq var' = d \wedge \dots$$

the sending event stores the data sent in a variable and then the receiving event is free to read it.

Modules declare constants, variables and inner Modules The variables are *public* in that they are visible to an outer Module. We can define the behaviour of a Module using a *Next* state predicate and an *Init* predicate

$$\text{Spec} \triangleq \text{Init} \wedge \text{Next}$$

Note the **Next** predicate is the critical component that defines what next step the Module can take given its current state.

The variables in a Module can be hidden, made *private* as we will see shortly.

A module **N** can declare a predefined module **Chan**

$$\text{InChan} \triangleq \text{INSTANCE Chan} \dots$$

An action **get** in **M** can include one of the actions in **InChan** simply by adding a conjunction the action **input** from the module **InChan** by:

$$\text{get} \triangleq \text{InChan!input}(\dots) \wedge \dots$$

13.1 Partial specifications

To create partial specifications that is to leave some thing unspecified define what is to be unspecified as a parameter. This is the same as indexed processes.

Let us assume that Module **Chan** contains a variable **q**. Now we can declare a parametrised Module **Chan(q)** where the parameter will instantiate **Chans** variable **q**

$$\text{Inner}(q) \triangleq \text{INSTANCE Chan} \dots$$

13.2 Variable hiding

Another use of parametrised Modules is when to hide a variable or put another way make a variable *private*. For example if we want the variable **q** of the inner module **Chan** to be *private* then we also declare

$$\text{Spec} \triangleq \exists q : \text{Chan!}(q)!\text{Spec}$$

Note this kind of variable hiding is not a mainstream idiom of TLA+ but is essential to our visual verification:

In fact, for most applications, there's no need to hide variables in the specication. [?, page 41]

If the declared modules state is treated as private, not accessed directly then the action **get** is a superposition refinement of the action **input**. Hence with some restrictions **INSTANCE** can be used to define superposition refinement between modules (Event-B).

Modules with private state are like processes, actions are like process events and a modules *Init* action can take the place of a processes start state.

13.3 Composition and Decomposition of specifications

Let a system be composed of two components, two processes or Modules. We might, quite reasonable want to reason in various related styles. We might start with an abstract system specification and want to decompose it into two separate parts alternatively we might go in the reverse direction and start with the specification of two components and want to compose them to form a system specification. Event B

provides a way to verify refinement steps having previously defined both the abstract system specification and the more concrete composition of two components. This technique is agnostic as the direction you are taking.

TLA+ formalises the behaviour of composition to two Modules as the disjunction of the behaviour of the component Modules. Slight variants are able to capture an interleaving or true concurrent behaviour.

13.4 Further Comparison with processes

Process algebras are built from a set of events and process operators.

A binary process operator is a bit like a parent module declaring as an instance two child modules. Such a declaration must give the parent access to the state of the child but the child modules will be private to the parent. In particular both the state and actions of the children can not be seen outside of the parent.

Process algebraic parallel composition is a binary process operator. Event synchronisation can be formalised by introducing an action that contains the conjunction of the two synchronised actions. Non synchronised child events are either lifted to the parent by introducing a parent action that simply includes the child action or can be blocked by not lifting them.

In TLA+ there is no counterpart to event hiding or testing equivalence. TLA+ makes use of temporal logic to specify the behaviour of a Module there is no ability to hide parts of an implementation or detailed specification to build a more abstract specification. Consequently if we were to add shared state to the processes our tool uses then the style of analysis would be quite different to that in TLA+.

14 Syntax

Processes are first defined then any you wish to view as an automata appear in the list see below:

```
const Max = 2
processes {
A = a ->STOP.
B = b -> STOP.
C = A || B }
automata B,C.
```

Event type	Symbol	Meaning
handshake	a	a synchronises with a both must be ready
non blocking send	a!	need not wait - can not be blocked
receive	a?	waits for a! synchronises to become a!

In above three processes are defined but only two displayed. The constant Max is bound to the value 2.

	atomic	indexed
Prefixing	A = act->P	if (i<N) then (act[i]->P[i+1]) else P[0]
		Money = C[1], C[i:1..N] = (when(i<N) coin->C[i+1] when(i==N) coin->C[1]).
Choice	A = a->P b->Q	Farmer = ([i:0..N].task ->W[i]), W[i:0..N] = ([i].end->Farmer).
Labeling	lab:P	see below
Parallel	A = (P Q)	Workers = (forall [i:0..N] ([i]:Worker)).
Relabeling	P/{new/old}	P/{new[i:0..N]/old[i]}
Hiding	P\{act}	P\{act[i:0..N]}
Keep	P{s} \triangleq P\{Act - s}	

For processing automata:

abstraction	abs(P)	fair removal of τ events
abstraction	abs{unfair}(P)	unfair removal of τ events
hiding	hide{S}(P)	\triangleq abs(P\{S})
simplification	simp(P)	for the simplification of automata
hiding index x	R = R1\${x}	builds symbolic automata

For processing automata use the following operations within:

operation \{ ... \}

Bisimilar	A ~ B	A and B are bisimilar
	A! ~ B	not bisimilar
Complete Tr	A#B	A and B are complete trace equivalent
	A!#B	not complete trace equivalent
Failure	A * B	A and B are failure equivalent
	A! * B	not failure equivalent
fair divergence		remove all τ loops
not fair divergence		replace τ loops with deadlock

References

- [BW90] J. C. M. Baeten and W. P. Weijland. *Process Algebra*. Cambridge Tracts in Theoretical Computer Science 18, 1990.
- [BW98] Ralph-Johan J. Back and J. Von Wright. *Refinement Calculus: A Systematic Introduction*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 1998.
- [dNH84] R. de Nicola and M Hennessy. Testing equivalences for processes. *Theoretical Computer Science*, 34, 84.
- [Hen88] M Hennessy. *Algebraic Theory of Processes*. The MIT Press, 1988.
- [Hoa85] C.A.R. Hoare. *Communicating Sequential Processes*. Prentice Hall International Series in Computer Science, 1985.
- [ReS04a] S. Reeves and D. Streader. Unifying state and process determinism. Technical report, University of Waikato, <http://hdl.handle.net/10289/1001>, 2004.
- [ReS04b] Steve Reeves and David Streader. Atomic Components. In Zhiming Liu and Keijiro Araki, editors, *Theoretical Aspects of Computing - ICTAC 2004: First International Colloquium*, volume 3407 of *Lecture Notes in Computer Science*, pages 128–139. Springer-Verlag, September 2004.
- [ReS09] Steve Reeves and David Streader. Guarded operations, refinement and simulation. In *Proc Fourteenth BAC-FACS Refinement Workshop (REFINE 2009)*, [doi:10.1016/j.entcs.2009.12.024](https://doi.org/10.1016/j.entcs.2009.12.024), volume 259 of *Electronic Notes in Theoretical Computer Science*, pages 177–191, Eindhoven, The Netherlands, 2009. Elsevier.
- [ReS11] Steve Reeves and David Streader. Contexts, refinement and determinism. *Science of Computer Programming*, DOI: [10.1016/j.scico.2010.11.011](https://doi.org/10.1016/j.scico.2010.11.011), 2010.
- [Tay99] P. Taylor. *Practical Foundations of Mathematics*. Cambridge University Press, 1999. Cambridge studies in advanced mathematics 59.
- [vG90] R. J. van Glabbeek. Linear Time-Branching Time Spectrum I. In *CONCUR '90 Theories of Concurrency: Unification and Extension*, LNCS 458, pages 278–297. Springer-Verlag, 1990.
- [vG93] Rob J. van Glabbeek. The Linear Time - Branching Time Spectrum II. In *International Conference on Concurrency Theory*, pages 66–81, 1993.

15 Tool Development

Currently toll builds finite state automata. Needs options to build symbolic automata, finite state Petri Nets and symbolic Petri Nets.

Quality Code The equation processor allows algebraic equational rules to be checked by building a set of test processes and then instantiating each variable to generate a ground equality that can be validated. We currently use this to generate thousands of tests. This has improved the code quality significantly.

1. *Algorithm design + documentation + Review*
2. *Code review prior to committing*
3. *Public Interfaces Java Docs + documentation*

1. Interface for Automata made public and documented.
2. Support for plug-in style development.

15.1 Technical overview

Ongoing review and discussion!

1. Asynchronous architecture with processing in a docker containers
2. Compile each process separately and only recompile if that part of the document has changed.
3. Isabelle is now designed for *proof as a service* with Isabelle running in docker containers in the background. Alas this appears to be work in progress and we might have to stick to Z3 at the moment.

15.2 TODO

1. Add
 - (a) co-events
 - (b) probabilistic choice (may well postpone till later)
2. ****start** Add Petri Nets,**
 - (a) compile to Petri Net Petri (and location) The tool takes as input process specifications and by default builds finite state automata as approximations. The function **Petri** can be applied to a process specification where it must build Nets for all processes below it in the **Process Hierarchy**. Define $\cdot \|_n$ on Petri Nets initially only for hand shake synchronisation.
 - (b) to define choice between parallel composition define the **Cross product of two sets of Places**. Consider example $(A \| B) | (P | Q)$
 - (c) Implement Token Rule to build automata **TokenRule**

Test using equalities: $\forall P : P \sim \text{TokenRule}(\text{Petri}(P))$

$\forall P, Q : P \parallel Q \sim \text{TokenRule}(\text{Petri}(P) \parallel_n \text{Petri}(Q))$

As a side effect this will speed up parallel composition.

3. Symbolic processes closely related to Petri Nets shown above :

- (a) Define $_ \parallel _$ on symbolic processes. Integrated with Petri Nets. To apply index hiding after abstraction need to track indexes
- (b) Define S2A a function that maps symbolic processes to atomic processes.
- (c) **Add Process invariants** for Z3
- (d) Use Isabelle in place of Z3 **but only if headless Isabelle ready**

Test like Petri Nets using equalities: $\forall P : P \sim \text{S2A}(P\$\{x\})$

$\forall P, Q : P \parallel Q \sim \text{S2A}(P\$\{x\} \parallel Q\$\{y\})$

4. Fast abstraction option for failure semantics relates abstraction in Petri Nets to abstraction in automata!

- (a) On automata
 - i. for n so that $(\forall t : t \in \text{pre}(n) \rightarrow \text{name}(t) = \tau) \wedge (\forall r \bullet r \in \text{post}(n) \rightarrow \text{name}(r) = \tau)$
 - ii. Replace $n, \text{pre}(n), \text{post}(n)$ with $\{x \xrightarrow{\tau} y \bullet x \in \text{pre}(\text{pre}(n)) \wedge y \in \text{post}(\text{post}(n))\}$
 - iii. better done in Labelled Petri Nets
- (b) On Petri Net: For transition t let $\text{post}(t) = t\bullet$ and $\text{pre}(t) = \bullet t$.
 For set of places R let $\text{pre}(R) = \{t : t\bullet \cap R \neq \emptyset\}$ and let $\text{post}(R) = \{t : \bullet t \cap R \neq \emptyset\}$
 - i. for transition r such that $\text{name}(r) = \tau$
 - ii. if $(\forall t \bullet t \in \text{pre}(r\bullet) \rightarrow \text{name}(t) = \tau) \wedge (\forall s \bullet s \in \text{post}(\bullet r) \rightarrow \text{name}(s) = \tau)$
 - iii. Replace $r\bullet, \text{pre}(r\bullet), \text{post}(r\bullet)$ with $\{\bullet x \xrightarrow{\tau} y\bullet : x \in \text{pre}(r\bullet), \wedge y \in \text{post}(r\bullet)\}$

Tested using $\text{abs}(P) =_F \text{fabs}(P)$

5. Add probabilities

- (a) probabilistic choice - reactive
- (b) abstraction will produce generative probability
- (c) could build lifting into abstraction - start set of nodes not just node?

6. Enhanced debugging:

- (a) Show shortest trace to expose inequality
- (b) Visualise bisimulation state equality between different automata or Petri Nets
- (c) keep appending to console output and add button to clear.
- (d) debugging tool tips when you hover over:
 - i. events \Rightarrow evaluation of variables in scope

ii. nodes \Rightarrow evaluation of variables in scope

7. **Compute Failure semantics** needed for example below.

(a) Apply subset construction to P to build a DFA, $nfa2dfa(P) = D$

$$\text{as } \forall n \in N_D : \exists ! S_D \xrightarrow{tr} n \text{ hence } \exists n2tr_D(.) : S_D \xrightarrow{n2tr_D(n)} n$$

(b) Annotate each DFA node $n \in N_D$ with Ready sets taken from nodes in P .

$$\{\pi(v) : v \in N_P \wedge S_P \xrightarrow{n2tr_D(n)} v\}$$

i. DFAR normal form add Ready set but keep A and drop B if $A \subseteq B$

(c) Compute DFAR equality

8. Isabelle extension overview:

- (a) Isar text managed by Jedit (JEdit is a plugin extensible editor designed for the construction of IDEs). Isabelle/Scala is used to manage to the JEdit
- (b) Could define functions, bisimulation, abstraction, automata building in Isabelle
- (c) Could keep the algorithms in Java and just use Isabelle to do symbolic to atomic conversion.

9. **Gui to define event renaming and synchronisation - needs Petri Nets**

10. **Syntax alternatives?** Variables are currently named and can be read but writing to variable is done by writing to the index at the same location at which the variable is declared.

Alternatively declare variables in a process and:

- (a) s and use a common variable assignment syntax $x := x+1$ or
- (b) use Z specification style primed variables for post state $x' = x+1$

11. **Add Process invariants**, need to look at TLA+ syntax / parsing + import libraries

12. Support for SDN may be TLA+

15.3 Link with TLA+

TLA+ has

1. Theories + data types
2. syntactic substitution
3. term simplification
4. expansion to finite state model.

Process Tool has

1. visualisation
2. parallel composition + synchronisation
3. abstraction
4. simplification bisim + failure + trace

Stepped development of a Bridge between TLA+ and PA.

Step Zero sequential processes with atomic events to TLA and back. Not sure if this should be to TLA text or parse tree!

1. **P2M**: $Process \rightarrow Module$ Process mapped to Module with additional State variable.
This will add the ability to check process satisfies temporal logic specification
2. **M2P**: $Module \rightarrow Process$ Module mapped to atomic automata using built in expansion

Add tests to test directory for any atomic processes P:

1. $M2P(P2M(P)) \sim P$

Tests added to repository can be refactored by adding new versions of P. (Might be worth defining algebra over process variables and set of processes to instantiate process variables)

Step One Define \parallel_M , parallel composition of TLA+ modules:

1. Component modules as INSTANCE declaration.
2. Event synchronisation defined as conjunction of component actions
3. default lifting of non synchronising actions.

Add tests for any atomic processes P and Q:

1. $M2P(P2M(P) \parallel_M P2M(Q)) \sim P \parallel Q$

Second step Change PA with indexed state and events to make use of TLA+ numbers

1. add import statement to PA
2. change PA to use TLA code for parsing and simplification of numbers:

Add tests using indexed processes

Third (2a) Code M2sP based on M2P but to return a symbolic numeric process:

1. **M2sP**: $Module \rightarrow Process$ Module mapped to symbolic automata.

For any symbolic processes P and Q use TESTS:

1. $M2sP(P2M(P)) \sim P$
2. $M2sP(P2M(P) \parallel P2M(Q)) \sim P \parallel Q$

Forth step use existing TLA code to add data theories to PA

1. extend PA to parse theories and display
2. extend M2sP and P2M

Fifth (3a) step add symbolic abstraction to PA

1. **implement** $sabs(-)$ symbolic abstraction using TLA+ code:
2. add tests
 $S2A\{x\}(sabs(P\{x\})) \sim abs(P)$

Sixth (3b) step add symbolic bisimulation PA

1. **implement** $ssimp(-)$ **symbolic bisimulation using TLA+ code:**
2. add tests

$$S2A\{x\}(ssimp(P\{x\})) \sim simp(P)$$

Bits and pieces.

1. Find TLA+ parse tree for module, TLA_{PT}
2. build a TLA_{PT} for a Petri Net
3. Find TLA+ term evaluation (tree for an action)
4. computation options:

(a) implement symbolic abstraction	(a) expand
(b) implement symbolic equality	(b) atomic abstraction
(c) expand	(c) atomic equality

15.4 Extensions TO DO

Below is a list of interesting projects that could be SWEN302/489 or even MSc they are given in particular order. Each project has interest both academically and pragmatically how easy they are to implement depends upon the state of the code base and no extension is of interest unless backed up by extensive executable tests.

1. Add support for Event B style reasoning.
2. Add probabilistic choice.
3. Add the ability to model check algebraic rules by generating "all" process models of a fixed size and verifying the rules against these processes.
4. Add interrupts. see after "`\end{documnet}`"
5. Add hierarchical processes (where the state of one process becomes a whole process) the result will be include adding signals that will better model interrupts.
6. Add event refinement
7. Add δ events - unobservable and blocked + model known
8. Code generate from the models
 - (a) The Go programming language has Go routines that can communicate via the CSP style event synchronisation that we use here. Consequently might be an easy target language for code generation.
 - (b) Java code - and apply the specification mining tool to rebuild the automata
 - (c) Ada code - and use SPARK Ada theorem prover to symbolically verify the indexed models for all values of the index

16 Bugs and examples

16.1 Bug 1 Abstraction (Solved)

16.2 Bug 2 Broadcast communication (Solved)

16.3 Bug 3 restricted rendering (Solved)

16.4 Need for failure semantics(Not a bug)

```
automata {
/* Question is 2 place buffer with: two faulty one place buffers
   = 2 place buffer with: one faulty and one good one place buffers?
*/
FA = (zero.in ->(zero.out->FA|drop->FA) | one.in ->(one.out->FA|drop->FA)).
FAs = simp(abs(FA\{drop})).
FB = (zero.in.[d:0..N] ->(zero.out.[d]->FB|drop->FB) |
      one.in.[e:0..N] ->(one.out.[e]->FB|drop->FB))\{drop}.
FBs = simp(abs(FB)).

Next = b:FAs/{a.zero.out/b.zero.in,a.one.out/b.one.in}.
First = a:FAs.
Two = (First|| Next)\{a.zero.out,a.one.out}.
TwoS = simp(abs(Two)).

Buff = a.one.in ->a.one.out->Buff | a.zero.in->a.zero.out->Buff.
One = (Buff|| Next)\{a.zero.out,a.one.out}.
OneS = simp(abs(One)).
}


operation {
One !~Two.
One # Two.
/* these are not bisimilar but are trace equivalent
Let One->a.one.in -> X1 and Two->a.one.in -> X2 where X1!=One and X2!=Two
Now X1 can not forget the input "one" where as X2 can.
But are they failure equivalent?
*/
}
```

17 Statefull syntax

1. Declare constants in each process
2. Declare variables in each process
3. Process = process name, indexed process or **assignment statements**
4. **guard event** \rightarrow **process** OR **event = guard process**

17.1 Translation:

```
CM = (teaButton->takeTea->STOP
    | coffeeButton->takeCoffee->STOP).
```



Introduce a variable *St* to distinguish different processes (nodes).

```
CM = const N;
var St:0..N;
init = St:=0;
events =
  when(St==0) teaButton->St:= 1
  |when(St==0) coffeeButton->St:= 2
  |when(St==1) takeTea->St:=3
  |when(St==2) takeCoffee->St:= 4.
```

```
CM = const N;
var cnt:0..N;
init = cnt:=1;
teaButton = when(St==0) St:=1
coffeeButton = when(St == 0) St:=2
takeTea = when(St == 1) St:=3
takeCoffee =when(St == 2) St:=4.
```

1. $C[1]$ into initialise first index to 1
2. $C[cnt:1..N] = \dots$ into declaration of *cnt* at first index and rhs as process definition
3. $C[cnt+1]$ into process after $cnt:=cnt+1$.

```
Money = const N;
var cnt:0..N;
init = cnt:=1;
events =
  when(cnt<N) coin -> cnt:=cnt+1
  |when(cnt==N) coin -> cnt:=1.
```

```
Money = const N;
var cnt:0..N;
init = cnt:=1;
coin = when(cnt<N) cnt:=cnt+1
      |when(cnt==N) cnt:=1.
```

18 Symbolic Processes **mostly 2b added 2017**

Simple process specifications, those with out indexes, are "atomic" specifications that produce atomic automata. An atomic automata has atomic transitions with a single atomic name and a finite set of atomic nodes each representing one state.

Indexed process specifications represent processes with variables, the indexes, and may be infinite state. The indexed specifications are expanded to a finite state approximation of the underlying and potentially infinite state automata. They do this using defined bounds such as `const N = 4`.

Alternatively we could define symbolic automata, automata with variables. The state of the process, represented by a symbolic automata, is the pair (n, μ) where μ is an evaluation i.e. a mapping from variables to values and n is a node of the automata.

$$\mu : Var \rightarrow Val$$

Hence the node of a symbolic transition only tells you part of the state of the process. When a event of a symbolic process is actually executed the state of the process both moves from the pre-node of the transition to the post-node of the transition and the evaluation of the variable changes as defined by the assignment. Symbolic execution represents a whole set of actual executions.

The symbolic transitions need to be annotated with its name plus a boolean guard and an assignment. Let a transition $t1 \triangleq (n1, g1, ev1, a1, m1)$ be represented as

$$n1 \xrightarrow{g1, ev1, a1} m1.$$

Both the guards $g1$ and assignments $a1$ may contain the process variables. Note symbolic Petri Nets can be built by adding the variables from a process to the places that represent the state of that process.

Our tool takes process specifications \mathcal{P} and, by default, builds finite state automata \mathcal{P} . But now we want to prevent the expansion of indexes and build symbolic automata from the specification $\mathcal{P}\{x\}$.

The execution of a guarded event can only occur when the guard is true and then the assignments of the transition are applied hence building another evaluation. Two essential functions needed in the definition of the application of assignments a , written $._@a$ are:

1. syntactic substitution and
2. simplification.

Simplification certainly needs to be out sourced to proof services as many decades of work has gone into developing such algorithms. And doing this allows for the development of theories of specific data types that include both their definition and the proof of rules of inference used in simplification. Essential this provides an extensible proof engine.

Syntactic substitution look easy but either you need to; hard bake in the language the data types used or out source both parsing and substitution. *Initially hard bake in integers, lists and sets?*

Reasoning about symbolic processes requires that we can concatenate sequential transitions. This can be achieved using two basic standard techniques, *symbolic execution* and backward reasoning via *Hoare Logic*. Hoare Logic tells us how to compute the weakest precondition prior to an assignment `ass` of a post condition `bg` and is written `bg@ass`. Whereas symbolic execution tell us how to combine the sequential execution of two assignments `a1` and then `a2` into a single assignment that is written `a1@a2`

Let assignment `ass` be a set of assignments all applied in parallel $\{x := E, y := F\}$ and let semicolon be used to compose assignments sequentially.

Compute the weakest precondition of an assignment for a known post condition:

$$\{P[E/x]\}x := E\{P\}$$

In our notation the precondition $P@\{x := E\} \triangleq P[E/x]$

Use symbolic execution to remove the sequencing of assignments.

$$x := F(x, y); y := G(x, y) = \{x := F(x, y), y := G(F(x, y), y)\}$$

In our notation $\{x := F(x, y)\} @ \{y := G(x, y)\} \triangleq \{x := F(x, y), y := G(F(x, y), y)\}$

19 Index freezing (z3 - headless Isabelle)


Closely related to Petri net Construction and Token Rule.

The default interpretation, semantics, of an indexed process definition is its finite state expansion using the declared values of the parameters. An alternative interpretation is to not expand some indexes but to leave them as unknown and define symbolic events.

```
S1 = X[0],
    X[x:0..N] = (when (x < N) out[x] -> X[x+1] |
                when (x == N) stop -> STOP).
```

S = S1\$ {x}.

1. assignment $x := x + 1$ is missing
2. presentation order *guard name assignment*
3. 2 should be N and \$x should be x



For the parallel composition of indexed processes we will need to α convert index names to prevent name clashes. This can be achieved by prefixing an index name with its process name, hence x above would become S1.x



```
R1 = input[y:0..N] -> X[y],
    X[x:0..N] = (when (x < N) out[x] -> X[x+1] |
                when (x == N) stop -> STOP).
```

R = R1\$ {x}.

T = R1\$ {x, y}.

S = R1\$ {y}.


1. S = R1\$ {y} currently fails to build.



In the above example both indexes x and y are needed.

```
Lift = L[0][0],
    L[x:0..N][to:0..N] = btnto[i:0..N] -> L[x][i] |
                        when (x != to) move -> L[to][to].
```

All three Indexes x, to and i are needed.



Need a symbolic expansion mapping, **S2A** that takes as input a symbolic automata and returns an atomic automata.

$$\begin{aligned} R &= R1\{x\}. \\ T &= S2A(R\{x\}). \end{aligned}$$


The result of expanding all variables should be the same as building the atomic automata from scratch, hence $R \sim T$ in above. (Note building from the automata/Petri Net allows interacting with the diagram)

19.1 Need for index freezing

Take a look at the examples below.

```
const N = 1
automata {
  Buf1data = input[i:1..N] -> F[i],
    F[d:1..N] = output[d] -> Buf1data.
  Buf2a = (a:Buf1data)/{b.move[x:1..N]/a.output[x]}.
  Buf2b = (b:Buf1data)/{b.move[x:1..N]/b.input[x]}.

  B2indexed = Buf2a||Buf2b.
  Bu2x = B2indexed${x}.
}
```



The two place buffer that, in each place can only hold the number 1.

What I do not understand is how **B2indexed** dose not have event **a.output[1]** but **Bu2x = B2indexed\${x}**. dose?

A more complex example is:


```

const N = 3
automata {
  Buf1data = input[i:1..N] -> F[i],
    F[d:1..N] = output[d] -> Buf1data.

  Buf2a = (a:Buf1data)/{b.move[x:1..N]/a.output[x]
  Buf2b = (b:Buf1data)/{b.move[x:1..N]/b.input[x]]

  Buf2data* = Buf2a|Buf2b.
  Buf2 = simp(abs(Buf2data\{b.move[x:1..N]})).
  //B2 = Buf2${x}.
  B2indexed = (Buf2a|Buf2b)\{b.move[q:1..N]}.
  Temp = B2indexed /{input/a.input[x:1..N],
    output/b.output[x:1..N]}.
  BestGuess = simp(abs(Temp)).
  //Bu2x = B2indexed${x}.
}
operation { Bu2x ~ BestGuess.}

```



The two place buffer that, in each place can hold a number from $\{1, 2, 3\}$ is defined and displayed. It is not as easy to see that it is a buffer as I wild like and if we could *freeze* the data in each place we might have simpler visual representation.

This will require some work to achieve. One tracking indexes so that we know what indexes exist in a process. For example what indexes are in Buf2 defined above? Hence what should we write

B2 = Buf2\${x}

19.2 Symbolic 2 Atomic mapping S2A

Symbolic execution is the obvious way to map S2A Symbolic Processes to Atomic Processes.

$$S2A(A\${x}) \sim A$$

State contains an evaluation of the variables.

All you need to do is build a new node for each distinct evaluation reached.

Add the set of initial evaluations to a "to do list" and repeatedly:

Remove the top node from the to do list and evaluate the boolean guard of each symbolic event and those that evaluate to true you apply the assignment to compute the new reachable state. If the new state already exists then add the transition ending at the corresponding node else add the state. When all events of the selected node have been processed either select the next node on the to do list else if the list is empty terminate.

Reasoning about symbolic systems is problematic as they are infinite state and consequently frequently require a degree of theorem proving. A vast amount of work has gone into both push button theorem proving and interactive theorem proving over the recent years. Yet push button theorem needs to be a lot stronger and interactive theorem proving a lot easier.

To code with data structures, such as lists and records, needs headless Isabelle. Initially we will only have Z3 and are restricted to integers. For data structures we need rules for rewriting, simplifying, the data structures. These rules constitute a *Theory* of the data structure and as they are heavily reused are defined and stored on file. Thus a procedure should import them not define them.

19.3 Symbolic parallel composition $- \parallel_s -$

The operations defined on atomic processes can be lifted to operations on symbolic processes. We first define $S2A$ that maps symbolic process to atomic processes and then lift the atomic operation Op_a to the symbolic operation Op_s so that:

$$S2A(Op_s(A, B)) \sim Op_a(S2A(A), S2A(B))$$

Symbolic parallel composition $- \parallel_s -$ is an extension of atomic parallel composition $- \parallel -$ where:

1. the nodes have a union of the indexes, suitably renamed to prevent name clashing
2. synchronising transitions have guards the conjunction of the component guards and assignments the union of the component assignments.

19.4 Symbolic abstraction

Let a transition $t1 \triangleq (n1, g1, ev1, a1, m1)$ be represented as $n1 \xrightarrow{g1, ev1, a1} m1$. We may refer to $n1$ as $t1_{pre}$, to $m1$ as $t1_{post}$ and refer to $e1$ as $t1_{en}$.

We need to compute $t1 : t2$ the transition representing the execution of $t1$ followed by $t2$. The execution of two transitions one after the other only occurs if the port node of the first transition is the pre-node of the second transition.

$$t1 : \tau : \frac{n \xrightarrow{g1, ev1, a1} m \quad m \xrightarrow{g2, \tau, a2} p}{n \xrightarrow{\underline{g1 \wedge g2 @ a1, ev1, a1 @ a2}} p} \quad \tau : t2 : \frac{n \xrightarrow{g1, \tau, a1} m \quad m \xrightarrow{g2, ev2, a2} p}{n \xrightarrow{\underline{g1 \wedge g2 @ a1, ev2, a1 @ a2}} p}$$

The guard $g2$ is the post condition to transition $t1$ hence to compute the weakest precondition we need to apply Hoare Logic thus compute $g2 @ a1$ which we add as a conjunction to the guard of $t1$. Thus if $t1_{pre} = t2_{post}$ we can add the transition :

$$t1 : t2 = (t1_{pre}, t1_g \wedge t2_g @ t1_a, t1_{en} : t2_{en}, t2_a @ t1_a, t2_{post})$$

but if $t1_g \wedge t2_g @ t1_a \neq \text{False}$ then this transition can never be executed and hence can be removed without effecting the behaviour of the process.

We can construct tests from:

$$S2A(abs(P)) \sim abs(S2A(P))$$

19.5 Symbolic simplification

We certainly want property:

$$S2A(simp(P)) \sim simp(S2A(P))$$


this could certainly be achieved if symbolic bisimulation simply matched transitions $t1$ and $t2$ by

$$t1_b \Leftrightarrow t2_b \wedge t1_{en} = t2_{en} \wedge t1_a \Leftrightarrow t2_a.$$

It might be possible to further simplify the symbolic automata while preserving the stated property.

20 Symbolic Petri Nets **2b added 2017**

The symbolic Petri Nets are an *orthogonal* combination of the symbolic extension to automata and the construction of Petri Nets rather than automata. Thus we have a square with finite automata, finite Petri Nets, symbolic automata and symbolic Petri Nets on the corners. Mappings between adjacent edges are functions that should preserve the semantics of the processes, be monotonic with respect to refinement and congruent with respect to the operators.



Index freezing prevents the building of finite state approximations of the index and instead builds a symbolic process. The construction $Symb(_)$ freezes all indexes but in the above square could be replaced by the more general index freezing $_ \{x, \dots\}$ of any set of indexes.

We have seen the partial expansion of both Petri Nets and of symbolic automata and clearly this can be generalised to a partial expansion of Symbolic Petri Nets.

20.1 Indexed Petri Nets

Indexed, specifications can be turned into FSPN in two quite separate ways. For a sequential indexed process $\mathcal{P}_{\mathcal{I}}$:

1. expand the index \mathcal{I} converting $\mathcal{P}_{\mathcal{I}}$ into a sequential FSPN
2. turn the index \mathcal{I} into a FSPN and the sequential process \mathcal{P} into another then build, $\mathcal{P} \parallel_n \mathcal{I}$, net parallel composition of the process and the index.

Both options can be applied to PTokenRule the partial expansion of a net into an automata. All options should result in interleaving equivalent Petri Nets.